

Z-test for Proportions in Statistics

A Z-test for proportion is a statistical hypothesis test used to determine whether the proportion of successes in a sample is significantly different from a hypothesized population proportion.

Note: It is particularly useful when dealing with categorical data and binary outcomes.

Why Use a Z-test for Proportions?

When we want to compare proportions (like the percentage of people who prefer one brand over another), we often use a statistical test. There are two common tests: the Z-test and the t-test.

The Z-test is like a shortcut that's really handy when we have a lot of data. It helps us figure out if there's a meaningful difference between proportions. We use it when we know or have a good guess about the standard deviation of the differences between proportions.

But sometimes, people use the t-test instead, even though it's not technically the right tool for the job. The t-test is meant for comparing averages, not proportions. However, when we have a lot of data, the results from a t-test end up looking pretty similar to what we'd get from a Z-test.

So, while the t-test isn't the best choice for proportions, it can still give us useful information, especially when we have a big bunch of data to work with.

▼ Two-Proportions Z-test

Compare the proportions p_1 and p_2 of two populations.

▼ Hypothesis

$$H_0 : p_1 = p_2$$

$$H_1 : p_1 \neq p_2$$

Under the null hypothesis, the two proportions are the same.

▼ Z-statistic

$$\hat{p}_1 \sim N(p_1, \frac{p_1(1-p_1)}{n_1})$$

$$\hat{p}_2 \sim N(p_2, \frac{p_2(1-p_2)}{n_2})$$

$$\hat{p}_1 - \hat{p}_2 \sim N(p_1 - p_2, \frac{\hat{p}_1(1-\hat{p}_1)}{n_1} + \frac{\hat{p}_2(1-\hat{p}_2)}{n_2})$$

Under H_0 , $Z \sim N(0, 1)$.

$$Z = \frac{\hat{p}_1 - \hat{p}_2}{\sqrt{\hat{p}(1-\hat{p})(\frac{1}{n_1} + \frac{1}{n_2})}} \text{ and } \hat{p} = \frac{k_1 + k_2}{n_1 + n_2} = \frac{n_1\hat{p}_1 + n_2\hat{p}_2}{n_1 + n_2}$$

This pooled estimate \hat{p} is similar to a weighted mean, but with two proportions.

A two proportion z test is conducted on two proportions to check if they are the same or not.

Question:

A company wants to improve the quality of products by reducing defects and monitoring the efficiency of assembly lines. In assembly line A, there were 18 defects reported out of 200 samples while in line B, 25 defects out of 600 samples were noted. Is there a difference in the procedures at a 0.05 alpha level?

Formulate Hypotheses:

- Null Hypothesis (H_0): There is no difference in the defect rates between assembly lines A and B ($p_A = p_B$)
- Alternative Hypothesis (H_1): There is a difference in the defect rates between assembly lines A and B ($p_A \neq p_B$).

Method 1: Using standard formula library

```
In [9]: import numpy as np
        from scipy.stats import norm
```

```
In [10]: # Sample sizes and defect counts for assembly lines A and B
        defects_A = 18
        samples_A = 200
        defects_B = 25
        samples_B = 600
```

```
In [11]: # Calculate sample proportions
        p_hat_A = defects_A / samples_A
        p_hat_B = defects_B / samples_B

        # Calculate pooled sample proportion
        p_hat_pooled = (defects_A + defects_B) / (samples_A + samples_B)

        # Calculate standard error
        SE = np.sqrt(p_hat_pooled * (1 - p_hat_pooled) * (1/samples_A + 1/samples_B))
```

```
In [17]: # Calculate test statistic
        Z = (p_hat_A - p_hat_B) / SE

        print("Test Statistic:", Z)
        print("P-value:", 2 * (1 - norm.cdf(np.abs(Z))))

        # Two-tailed test, so multiply alpha by 2
        alpha = 0.05
        alpha *= 2

        # Determine critical value
        critical_value = norm.ppf(1 - alpha/2)
        print("Critical Value:", critical_value)
```

```
test_statistic: 2.624824049200042
P-value: 0.008669375420073067
Critical_Value: 1.6448536269514722
```

As $2.62 > 1.64$ thus, the null hypothesis is rejected and it is concluded that there is a significant difference between the two lines.

```
In [13]: # Make decision
if np.abs(Z) > critical_value:
    print("Reject the null hypothesis: There is a significant difference in defect rates")
else:
    print("Fail to reject the null hypothesis: There is no significant difference in defe
```

Reject the null hypothesis: There is a significant difference in defect rates between assembly lines A and B.

-- This code calculates the test statistic (Z) and compares it with the critical value for a two-tailed test at a significance level of 0.05.

- If the absolute value of the test statistic exceeds the critical value, we reject the null hypothesis and conclude that there is a significant difference in defect rates between assembly lines A and B.
- Otherwise, we fail to reject the null hypothesis.

▼ Confidence interval for the difference between 2 proportions

Typically, we used **unpooled** proportions instead of pooled estimate of proportions.

A level $(1 - \alpha) * 100\%$ confidence interval of $p_1 - p_2$

$$\hat{p}_1 - \hat{p}_2 \pm z_{\alpha/2} \sqrt{\frac{\hat{p}_1(1 - \hat{p}_1)}{n_1} + \frac{\hat{p}_2(1 - \hat{p}_2)}{n_2}}$$

```
In [28]: # Calculate confidence interval for the difference in proportions
margin_of_error = critical_value * SE
confidence_interval = (p_hat_A - p_hat_B - margin_of_error, p_hat_A - p_hat_B + margin_of_error)
print("Confidence Interval for the Difference in Proportions:", confidence_interval)
```

Confidence Interval for the Difference in Proportions: (0.023541532041150864, 0.0731251346255158)

Method 2: Using statsmodels library

```
In [14]: from statsmodels.stats.proportion import proportions_ztest
```

```
In [16]: # Perform z-test for proportions

# Counts of successes (defects) in the two samples
successes = np.array([defects_A, defects_B])

# Sample sizes of the two samples
sample_sizes = np.array([samples_A, samples_B])

# Perform two sample z test for proportions
```

```
# Perform two-sample z-test for proportions
stat, p_value = proportions_ztest(successes, sample_sizes, alternative='two-sided')

# Print the test statistic and p-value
print("Test Statistic:", stat)
print("P-value:", p_value)

# Set significance level (alpha)
alpha = 0.05

# Compare p-value with alpha
if p_value < alpha:
    print("Reject the null hypothesis: There is a significant difference in defect rates")
else:
    print("Fail to reject the null hypothesis: There is no significant difference in defect rates")

print(f"P-value: {p_value:.3f}")
```

Test Statistic: 2.624824049200042

P-value: 0.008669375420073044

Reject the null hypothesis: There is a significant difference in defect rates between assembly lines A and B.

P-value: 0.009

-- This code performs a two-sample z-test for proportions using the `proportions_ztest` function from the `statsmodels` library.

It then compares the obtained p-value with the significance level (alpha) to make a decision regarding the null hypothesis.

- If the p-value is less than alpha, the null hypothesis is rejected, indicating a significant difference in defect rates between assembly lines A and B.
- Otherwise, the null hypothesis is not rejected.

Two-Proportion Z-test:

Advantages:

- Allows for comparing two proportions from independent samples, making it suitable for evaluating differences between two groups. Suitable for large sample sizes, ensuring reliable results and robustness.

Disadvantages:

- Assumes that the sample proportions follow a normal distribution, which may not hold true for small sample sizes or when the population proportions are close to 0 or 1.
- Requires independence between samples and random sampling to ensure validity of results. Limited to comparing proportions from two groups, making it less suitable for comparisons involving multiple groups or factors.

▼ One-Proportion Z-test

Compare a proportion of a population to a constant.

Let p be the success rate of a large number n of independent Bernoulli trials.

Let \hat{p} be the observed success rate, that is the number of observed successes over the total number of trials.

$$n\hat{p} \sim \text{Binomial}(n, p) \sim N(np, npq)$$

$$\hat{p} \sim N\left(p, \frac{p(1-p)}{n}\right)$$

▼ Hypothesis

$$H_0 : p = p_0$$

$$H_1 : p \neq p_0$$

▼ Z-statistic

Under H_0 , $Z \sim N(0, 1)$.

$$Z = \frac{\hat{p} - p_0}{\sqrt{\frac{p_0(1-p_0)}{n}}}$$

A one proportion z test is used when there are two groups and compares the value of an observed proportion to a theoretical one.

The null hypothesis is that the two proportions are the same while the alternative hypothesis is that they are not the same.

Question:

Suppose an algorithm for Ad selection is implemented, and we aim to estimate the click-through rate (p) of users on the Ads selected by this algorithm. With access to data from 1000 users, the observed click-through rate is $p_{\text{hat}}=0.2$. We set a significance level of 5%.

Formulate Hypotheses:

- Null Hypothesis (H_0): The click-through rate (p) is equal to 0.15.
- Alternative Hypothesis (H_1): The click-through rate (p) is not equal to 0.15.

```
In [23]: # Given data
n = 1000 # Total number of users
p_hat = 0.2 # Observed click-through rate
p0 = 0.15 #Hypothesized click-through rate (under H0)
alpha = 0.05 # Significance level
```

```
In [24]: # Calculate the standard error
SE = np.sqrt((p_hat * (1 - p_hat)) / n)

# Calculate the z-score
Z = (p_hat - p0) / SE
print("Test Statistic (Z-score):", Z)
```

Test Statistic (Z-score): 3.952847075210475

```
In [25]: # Calculate the critical value
critical_value = norm.ppf(1 - alpha/2)
```

```
critical_value = norm.ppf(1 - alpha/2)
print("Critical Value:", critical_value)

# Calculate the p-value
p_value = 2 * (1 - norm.cdf(np.abs(Z)))
print("P-value:", p_value)

# Make a decision based on the p-value
if p_value < alpha:
    print("Reject the null hypothesis: There is evidence that the click-through rate is d
else:
    print("Fail to reject the null hypothesis: There is no evidence that the click-throug
```

Critical Value: 1.959963984540054

P-value: 7.72267955053696e-05

Reject the null hypothesis: There is evidence that the click-through rate is different from 0.5.

▼ Confidence interval for a proportion

A level $(1 - \alpha) * 100\%$ confidence interval for p

$$\hat{p} \pm z_{\alpha/2} \sqrt{\frac{\hat{p}(1 - \hat{p})}{n}}$$

To calculate the confidence interval for the click-through rate, we can use the formula:

$$\text{Confidence interval} = \hat{p} \pm Z \times SE$$

Where:

- \hat{p} is the observed click-through rate,
- Z is the critical value from the standard normal distribution corresponding to the desired confidence level,
- SE is the standard error of the proportion.

In [26]:

```
# Calculate confidence interval
margin_of_error = critical_value * SE
lower_bound = p_hat - margin_of_error
upper_bound = p_hat + margin_of_error

print("Confidence Interval: [{:.4f}, {:.4f}].format(lower_bound, upper_bound))
```

Confidence Interval: [0.1752, 0.2248]

One-Proportion Z-test:

Advantages:

- Provides a straightforward hypothesis test to determine if the observed proportion significantly differs from a specified value.
- Applicable to large sample sizes, ensuring robustness and reliability of results.

Disadvantages:

- Limited to comparing a single proportion to a specified value, making it less flexible for more complex comparisons.